

# Benchmarking freeRTOS on PC & Raspberry Pi

Stylianos Tsagkarakis  
MIEEC - FEUP  
up201911231@fe.up.pt

Alena Tesařová  
MIEEC - FEUP  
up201911219@fe.up.pt

**Abstract**—Project developed within the scope of the Embedded Systems curricular unit, with the objective of benchmarking FreeRTOS on the platform of Raspberry Pi 2B and Linux Kernel.  
**Index Terms**—Real-Time Systems, RTOS, Performance

## I. INTRODUCTION

An operating system is said to be real time when it schedules the execution of programs in time, handles system resources and gives a reliable basis for the development of software code. [1]

Timing is the most important factor in any real time applications. In real-time systems, all real-time tasks are differentiated based on their timing, such as sporadic, response time, deadline etc. Real time systems are classified in to two types' hard real-time systems and soft real time systems. Hard real time system means it should complete the task with in the deadline period otherwise its computation is useless. The damages caused by the hard real time systems are irreparable. Soft real time systems require performance assurances from the operating system.

## II. HARDWARE

### A. Specifications of hardware used for Benchmarks

- Raspberry Pi Raspberry Pi 2 Model B (ARM Cortex-A7 900MHz), 1GB RAM
- Ubuntu 18.04, Linux Kernel 5.3.0-53-generic, Intel(R) Core(TM) i7-6500U CPU @ 2.50GHz, 12GB RAM (1 core & 2 threads were used)

## III. EVALUATING RESOURCES

By reading documentation [5], related previous work [3] we managed to generate the array below stating characteristics of freeRTOS:

Multitasking	RR Scheduling	Priority
Preemptive or Cooperative	Yes	Unlimited
# Tasks	Compilers Supported	Kernel size (KB)
Unlimited	IAR/Keil/gcc	ROM: 2.7-3.6 RAM: 0.19

In order to complete this project we used the scheduler freeRTOS has already implemented. This scheduler works based on this policy: **Highest priority task is granted CPU time**. If multiple tasks have equal priority, it uses round-robin scheduling among them. Lower priority tasks must wait. It is important that high priority tasks don't execute 100% of the

time, because lower priority tasks would never get CPU time (starvation problem). It's a fundamental problem of real-time programming.

## IV. TESTING

### A. Timing

To count the desired value two timers were used that counted the time based on this assembly instruction:

```
asm volatile ("mrc p15,0,\%0,c9,c13,0": "=r" (cc))
```

and we could get a low level timer using the processor clock speed. By using software times to count values of nanosecond accuracy we would not have the same results.

### B. Tests

In this section we will briefly describe the tests and obtained results. In each test we measured the the time of one specific operation for at least 500 iterations to get better average value. Then, we run the test more times to really get reliable results. We performed 4 types of tests regarding:

- Cooperative context switch (Round Robin)
- Message Queue
- Semaphores
- Mutexes

## V. RESULTS

In cooperative scheduling we had 2 tasks with the same priority that were switching the context in round robin fashion. As we can see from the Table V, we got a median time 168 ns – so we can use it as a starting time for other tests using context switch. However, it is not the best time we got. The lowest time we got for Receive test (137 ns) – that is the time it takes to a task to receive a message from a message queue knowing that the message is ready for the task in a message queue (without blocking). This scenario requires only one task and that is why there is no context switch. Send scenario is very similar. In this case we count the time to send a message to a message queue. It takes approximately 182 ns so there is a 50 ns difference between send and receive. There can be for 2 reasons. The first is that every time we send a value to a message queue it makes a copy (not reference) of the message and the second is that if the queue is full it waits some time and then it rewrites the previous value in the queue.

While measuring message queues we also added 3 complex tests. The signal unblock scenario can be see in Figure 1. In this test we were measuring the context switch with sending a message to a message queue. The median time reached 528 ns, if we subtract the context switch time we get 360 ns. That is




Fig. 1. Send unblock

the corresponding time to send the message to the queue. The value is greater than a pure send because in this case we must include the blocking time because the queue might be full. By looking at the receive block, the situation is similar but with receive and we get the median of 863 ns. That is because of the blocking time as it tries to read the value for some time if there is no value.

The last complex test with message queues is a workload test. In this scenario we had also sender, receiver and another tasks with lower priority that were doing some work. Here, we were measuring the time to send the message but at the beginning of sender activity we set a delay so the lower priority task could be executed. The interesting part is in Figure 2 where we can see in what percentage of our measurement is translated to context switch. It is both expected and obvious that in the maximum values, workload would require more time than context switching. In general we achieved much better results, compared to previous work [4], [6], needing almost 10% less cycles throughout the test.




Fig. 2. Message passing - workload analysis

We can observe similar behavior with other tests also *e.g. semaphores, mutexes, etc.*

Regarding semaphores we can see that results are coherent with message queue results, which is expected since both tests operate under the same principle. As we can see freeRTOS takes twice as much cycles to take and block a semaphore than to signal and switch a new task.

In mutexes tests we can notice the average is very close to the minimum values. This shows us that maximum values are not expected and are probably caused rarely by "extreme" cases.

Comparing mutexes and semaphores, sem signal & mutex acquisition and sem wait & mutex release we can see that mutexes are slower which is rational. This is due to the many security checks performed by mutexes before locking or unlocking.

Name	Min [ns]	Median [ns]	Max [ns]
mq context switch	157	168	480
mq send	168	182	1719
mq receive	137	137	342,5
mq send unblock	512	528	563
mq receive block	832	863	1337
mq workload	535	544	2526
mutex release unblock	518	529	1009
mutex request block	1032	1056	1496
mutex pip	1026	1042	1495
mutex workload	528	529	780
mutex acquisition	119	132	943
mutex release	157	160	526
sem signal unblock	420	427	448
sem wait block	848	862	1393
sem signaling with prio	214	229	606
sem signal	121	125	1352,5
sem wait	100	100	289,5
sem workload	698	698	2335

TABLE I  
RESULTS MEASURED ON RASPBERRY

## VI. CONCLUSION

The advanced real-time systems will possess capabilities for high-speed data processing and communication which will require very high time bound processing than what is available in state-of the-art systems. This necessitates the need of improving the real-time mechanisms in RTOS for our future need. As proved in this report, freeRTOS might have room for improvement. To cope with these challenges, very high performance is necessary at all levels of implementation application level and system level. The results show that real time operating system have better response time compared to the general purpose operating system. In this paper we reviewed several Benchmarking techniques for measuring the performance of RTOS using real-time parameters. It is hoped that by providing insights into the Benchmarking techniques, this paper would help the researches in addressing the implementation challenges of RTOS for efficient realtime systems of tomorrow.

### A. Video

Short video presentation of our work: <https://www.youtube.com/watch?v=iywyAUN0roM>

### B. Members Contribution

Each member contributed equally into the making of this report. (50% - 50%)

## REFERENCES

- [1] Wei-Tsun Sun; Zoran Salcic; , "Modeling RTOS for Reactive Embedded Systems," VLSI Design, 2007. doi:10.1109/VLSID.2007.111
- [2] Yodaiken, Victor & Barabanov, Michael. (2000). A real-time Linux.
- [3] Weideman, N.H., Kamenoff, N.I. Hartstone Uniprocessor Benchmark: Definitions and experiments for real-time systems. Real-Time Syst 4, 353–382 (1992).
- [4] RTOS Benchmarking tests, Guillaume Champagne, <https://github.com/gchamp20/RTOSBench>.
- [5] FreeRTOS Scheduling, <https://www.freertos.org/implementation/a00005.html>.
- [6] Benchmarking Real Time Operating Systems, Guillaume Champagne, Michel Dagenais, May 6, 2019, <https://amdls.dorsal.polymtl.ca/system/files/RTOS%20%20Benchmarking.pdf>